

Introduction to R

Valentina Zangirolami

University of Milano-Bicocca
Department of Economics, Management and Statistics

May 12, 2025



Contact: valentina.zangirolami@unimib.it

- To install R: <https://cran.r-project.org/>
- To install R studio (R is required): <https://posit.co/downloads/>
- Documentation: <https://cran.r-project.org/manuals.html>
- Cheat sheets: <https://support--rstudio-com.netlify.app/resources/cheatsheets/>
- Book:
 - *Venables, W. N., Smith D. M. & the R Core Team (2021). An Introduction to R.*
 - Efficient R programming:
<https://csgillespie.github.io/efficientR/index.html>
- To find specific topics: <https://www.bigbookofr.com/>

- `help.start()`: shows the main links for programming in R
- `help()`: to obtain the description of a specific function (e.g. `help(mean)` or `?mean` shows the description of the function "mean")
- `help.search()`: to search a string on function's documentation

Basic commands - Workspace

- `getwd()`: to visualize the working directory
- `setwd(" mydirectory")`: to set a new directory
- `ls()`: to visualize all the objects saved in R space
- `rm(myobject)`: to remove myobject from R space
- `rm(list=ls())`: to remove all files from R space

Basic commands - Libraries

- `library()`: list of packages available
- `library(name package)`: to load a package
- `install.packages("name package")`: to install a package
- `update.packages(ask=FALSE)`: to update all packages
- `data(package='name package')`: to show the list of datasets available on the "name package"
- `data()`: to show the list of datasets available on all loaded packages
- `data(name dataset)`: to load a dataset from a loaded package

- lowercase and uppercase are different on R
- '#' can be used to make comments
- **Symbols for maths operation:** + (addition), - (difference), * (product with numbers), / (division with numbers), %/% (integer division), ** (exponentiation), %% (division remainder)
- **Logic operators:** != (not equal), & (AND), | (OR)

In R, there are several kind of data types:

- **Numeric:** real numbers with or without decimal points (e.g. 2.5, 2, ...)
- **Integer:** real numbers without decimal points. To specify an integer number we need to use the suffix L (e.g. `integer_number = 18L`)
- **Logical:** boolean values (TRUE and FALSE)
- **Complex:** complex numbers whose imaginary values are represented by the suffix i (e.g. 2i)
- **Character:** string values (e.g. `fruit = "Cherry"`)
- **Raw:** values as raw bytes

Special data types

R involves two special data types: **NULL** and **NA**. When assigning **NULL** to a variable, we convert it into an **empty object**. Instead, **NA** means non-available entry and usually represents a **missing value** in data analysis

```
1 > null_object <- NULL
2 > missing <- NA
3
4 > null_object
5 NULL
6 > missing
7 [1] NA
8
9 > print(1 + missing)
10 [1] NA
```


Variables and Basic Arithmetic

R can be used as calculator to evaluate simple expression.

Some example:

```
1 # Sum of two numbers
2 > sum_n <- 2 + 2.5 #NB: the decimal separator is .
3 > cat("The sum of the two numbers is", sum_n)
4 The sum of the two numbers is 4.5
```

```
1 # Division between two numbers with scientific notation
2 > x <- 1
3 > y <- 2.73e4
4 > print(x/y)
5 [1] 3.663004e-05
```

Math functions

The most common math functions are already implemented in R:

```
1 # Exp and log functions
2 > z <- exp(3)
3 > y <- log(1)
4 > print(z - y)
5 [1] 20.08554
6 # Pi greco
7 > pi
8 [1] 3.141593
9 # Trigonometric functions
10 > sin(0); cos(pi)
11 [1] 0
12 [1] -1
```

If we want to calculate $\log_2(4)$:

```
1 > log(x=4, base=2)
2 [1] 2
```

It's your turn!

1. Given a circle with a radius equal to 4. Calculate perimeter and area of the circle using R
2. Install the library dplyr and show the datasets available.
3. Try to read the help page of `factorial` and `choose`.
4. In how many ways can you arrange 7 out of 20 books on a bookshelf? Show the result in R. (Hint: `ex.3` can be useful)
5. Round the number $n=3.78957$ to two decimal places. (Hint: use `round` function looking the documentation)
6. Visualize all objects saved in R space and then remove them.

Vectors

A vector is an array of elements having the same data type. In R, to create a vector we can use `c()`:

```
1 > # Different ways to assign x to a vector
2 > x <- c(23.4, 6, 7.1, 6.9, 34.7)
3 > assign("x", c(23.4, 6, 7.1, 6.9, 34.7))
4 > x = c(23.4, 6, 7.1, 6.9, 34.7)
```

The length of the vector we just created can be retrieved by

```
1 > length(x)
2 [1] 5
```

The "`[...]`" operator can be used to draw out parts of a vector by specifying one or more indices of numeric type. **In R, the indices of a vector are defined starting with the value 1.**

```
1 > x[1]; x[c(2,3)]; x[c(1:3)]
2 [1] 23.4
3 [1] 6.0 7.1
4 [1] 23.4 6.0 7.1
```

If you want to create a sequence of values by following a specific condition, you can use `seq()` or `rep()`

```
1 > n<-10
2 > seq(1, n); seq(1, n, by = 2); seq(1, n, length.out = 5)
3 [1] 1 2 3 4 5 6 7 8 9 10
4 [1] 1 3 5 7 9
5 [1] 1.00 3.25 5.50 7.75 10.00
6 > rep(n, times=5); rep(n, each=5) # same
7 [1] 10 10 10 10 10
8 [1] 10 10 10 10 10
9 > rep(1:4, each=5); rep(1:4, times=5) # different
10 [1] 1 1 1 1 1 2 2 2 2 2 3 3 3 3 3 4 4 4 4 4
11 [1] 1 2 3 4 1 2 3 4 1 2 3 4 1 2 3 4 1 2 3 4
```

Math operations and useful functions

There are some functions to tackle vectors or do math operations

```
1 > x<-c(x, -6.8) # add a new value to the previous vector
2 > sort(x) # default: decreasing=F, i.e. ascending order
3 [1] -6.8 6.0 6.9 7.1 23.4 34.7
4 > floor(x) # extract the integer part of a decimal number
5 [1] 23 6 7 6 34 -7
6 > sign(x) # sign: -1 = negative, 0 = 0, 1 = positive
7 [1] 1 1 1 1 1 -1
8 > any(x>4) # is there at least a value in our vector >4?
9 [1] TRUE
```

- `abs()`, `sum()`, `mean()`, `sd()`, `median()`, `prod()`, `cumsum()` compute absolute value, sum, average, standard deviation, median, product and cumulative sum
- `min()`, `max()` and `range()` show min and max of a vector
- `head()` and `tail()` extract the head and the tail of a vector (6 elements)
- `summary()` shows summary statistics

Vectors of characters

The R language implements a data structure called "factor" that facilitates the analysis of categorical data. Factors are particularly useful for representing variables that can take only one category among a known and finite set of values (e.g., gender, type of employment, month of birth, ...). Values can have an ordering (e.g. hierarchies in the military) or not (e.g. hair color).

```
1 > days_list <- c("Monday", "Tuesday", "Wednesday", "Thursday",  
  ", "Friday", "Saturday", "Sunday")  
2 > factor(c("Monday", "Sunday"), levels = days_list)  
3 [1] Monday      Sunday  
4 7 Levels: Monday Tuesday Wednesday Thursday ... Sunday
```

The `table()` function is used to generate a table of absolute frequencies

```
1 > table(c("A", "B", "B"))  
2 data  
3 A B  
4 1 2
```

Matrix

Matrices are a generalization of vectors in two dimensions. As the vectors, all elements of a matrix must be of the same data type.

```
1 > matrix(1:10, nrow = 2, ncol = 5)
2      [,1] [,2] [,3] [,4] [,5]
3 [1,]    1    3    5    7    9
4 [2,]    2    4    6    8   10
5 > matrix(letters[1:10], nrow = 2, ncol = 5, byrow = TRUE)
6      [,1] [,2] [,3] [,4] [,5]
7 [1,] "a"  "b"  "c"  "d"  "e"
8 [2,] "f"  "g"  "h"  "i"  "j"
```

The `dim()` can be used to show the dimensions of a matrix

```
1 > x <- matrix(1:10, nrow = 2, ncol = 5)
2 > dim(x) # dimensions of the matrix x
3 [1] 2 5
4 > nrow(x); ncol(x) # total number of rows and columns
5 [1] 2
6 [1] 5
```


Math operations

- Matrix multiplication: `%*%`
- `t()` and `det()` to compute matrix transpose and the determinant

```
1 > x <- matrix(1:4, nrow=2); y <- matrix(c(5,10,33,2), ncol
   =2)
2 > x; solve(x) # solve the equation a %*% m = b for m. In our
   case, a=x and b=diag(2) (identity matrix)
3      [,1] [,2]      [,1] [,2]
4 [1,]     1     3 [1,]    -2  1.5
5 [2,]     2     4 [2,]     1 -0.5
6 > crossprod(x,y) # t(x) %*% y
7      [,1] [,2]
8 [1,]    25    37
9 [2,]    55   107
10 > solve(x,x) # solve(x) %*% x
11      [,1] [,2]
12 [1,]     1     0
13 [2,]     0     1
```

Math operations

- `eigen()`, `svd()`, and `qr()` can be used to compute eigenvalues and eigenvectors, singular value decomposition and QR decomposition
- `cbind()` and `rbind()` mean column-binding and row-binding. Let's see an example

```
1 > cbind(x,y) # join matrices by column
2           [,1] [,2] [,3] [,4]
3 [1,]      1   3   5  33
4 [2,]      2   4  10   2
```

It's your turn!

7. Calculate absolute and relative frequencies of vector ("Cat", "Dog", "Bird", "Bird", "Cat", "Cat", "Dog"). Print the maximum absolute frequency and the correspondent animal. (Hint: use `which()` function)
8. Define a vector called `z` by "pasting" the vectors `x`=("Red", "Green", "Yellow") and `y`=("Tomato", "Grass", "Sun") and select only those elements of `z` that begin with "r" and end in "o". (Hint: use `paste()` and `startsWith()` functions)
9. Create a matrix `A` containing the integers from 10 to 24 having 5 rows and 3 columns where `A` must be filled by column. Show the firsts 3 rows of the firsts two columns. Calculate the sum of `A` by row and column (Hint: `rowSums()`, `colSums()`). Calculate the transpose of `A`. Define a new matrix `B` of appropriate size to calculate the `AB` product (and calculate that product). Calculate the standard deviation of `B` by row and then by column (Hint: `apply()`)

Lists can contain elements with different data type. They are also called "recursive" structures in that one list can also contain another list.

```
1 > list(FALSE, 0L, pi, "ABC")
2 [[1]]      [[2]]
3 [1] FALSE   [1] 0
4 [[3]]      [[4]]
5 [1] 3.141593 [1] "ABC"
6 > x <- list(1, "A", matrix(1:4, 2, 2))
7 > x[2] # to extract a list
8 [[1]]
9 [1] "A"
10 > x[[2]] # to extract the second element
11 [1] "A"
```

- If you want to transform a list to a vector, you can use `unlist()`

10. Given a matrix X 2×2 containing the integers from 1 to 4. Try to apply `svd()` function to X and then reconstruct the input matrix X with the outputs of `svd`.

Probability functions

R software implements a set of functions having a common syntax to handle random variables (e.g. `rnorm`, `qnorm`, `dnorm`, `pnorm`).

The first letter identifies the objective of the function:

- **p** means "probability", the cumulative distribution function (c.d.f.)
- **d** means "density", the density function (p.f. or p.d.f.)
- **r** means "random", to generate random numbers
- **q** means "quantile", the inverse c.d.f.

The second part describes the random variable (e.g. `norm` for normal distribution, `pois` for poisson, `gamma` for gamma, `unif` for uniform).

```
1 > rnorm(n = 5, mean = 0, sd = 1) # vector of 5 normally
    distributed random numbers
2 > pnorm(0, mean = 0, sd = 1) # P(X<=0) where X ~ N(0,1)
```

11. Let $X \sim N(5, 5)$. Calculate the 0.7 quantile of X .
12. Let $Y \sim \text{Gamma}(1, 2)$. Simulate the extraction of $n=500$ random numbers from Y and save them into a object called "data". Show the summary statistics about it. (To guarantee the reproducibility you need to use `set.seed()`)

There are several kind of control structures in R:

- `if/else`: to test a condition
- `for`: to execute a loop a fixed number of times
- `while`: to execute a loop while a condition is true
- `repeat`: to execute an infinite loop
- `break`: to break the execution of a loop
- `next`: to skip an iteration of a loop

The most used control structures are `if/else` and `for`.

There are different ways to implement if/else conditions:

```
1 # One condition
2 if (TEST) {
3     do - something}
4 # Two conditions
5 if (TEST) {
6     do - something
7 } else {
8     do - something - else}
9 # More conditions
10 if (TEST) {
11     do - something
12 } else if (ANOTHER TEST) {
13     do - something - else
14 } else {
15     do - something - else}
```

For

A for loop can be implemented using the following

```
1 for (sequence of indices) {  
2   body of the loop}
```

Let's see an example: Generate 10 random numbers $X \sim \text{Poisson}(10)$ and divided them into even and odd vectors.

```
1 > set.seed(123)  
2 > x <- rpois(10, lambda = 10) # data generation  
3 > x <- sort(x)  
4 > even_x <- c(); odd_x <- c() # initialize vectors  
5 > for (i in seq_along(x)) { # seq_along => 1:length(10)  
6 +   if(x[i]%%2==0){  
7 +     even_x <- c(even_x, x[i]) # add value x[i] to vector  
8 +     even_x  
9 +   } else {  
10 +     odd_x <- c(odd_x, x[i])  
11 +   }
```

Functions in R

In R to create your own functions, you can call `function()` and assign it to an object.

```
1 > fahrenheit_to_celsius <- function(temp_F) {  
2 +   temp_C <- (temp_F - 32) * 5 / 9  
3 +   return(temp_C)}  
4 > fahrenheit_to_celsius(55) # to show the result  
5 [1] 12.77778
```

In R, it's not necessary to specify the return statement

```
1 > fahrenheit_to_celsius <- function(temp_F) {  
2 +   (temp_F - 32) * 5 / 9} # as above  
3 > fahrenheit_to_celsius(55)  
4 [1] 12.77778
```

It's your turn!

13. Try to define a function `f` which, given a numerical vector as input, returns a list containing the mean, variance, minimum and maximum of that vector. Assign an appropriate name to the elements of the list.
14. Write an R function, `num_seq(n)`, that returns a vector that has the form $\{2^2 \cdot 3, 3^2 \cdot 5, 4^2 \cdot 7, 5^2 \cdot 9, \dots\}$ and has exactly `n` terms.
15. Assume that a data set is stored in the vector `data` in R. Write a function `skewness(data)` that calculates the skewness index ($I = \frac{3(\bar{x} - Q_2)}{s}$) of the data set and returns either "The data is significantly skewed" or "The data is not significantly skewed" as appropriate.

Dataframe

R uses an object called "data.frame" (i.e., class "data.frame") to represent matrix of data. The rows of a data.frame correspond to statistical units (i.e., observations) and columns to variables.

```
1 > x <- data.frame(hair_color = c("Blonde", "Blonde", "Black",  
    , "Brown"), gender = c("F", "M", "M", "F"))  
2 > x  
3   hair_color gender  
4 1    Blonde      F  
5 2    Blonde      M  
6 3    Black      M  
7 4    Brown      F
```

For external datasets, you can use two functions: `read.table()` and `read.csv()`.

Apply, sapply, tapply, lapply

Instead of for loop, in R you can use `apply`, `sapply`, `tapply` or `lapply`. This family of functions allows us to apply a certain function to a certain data frame, list or vector and return the result as a list or vector depending on the function we use.

- `apply()`: to apply a function to the rows or columns of a matrix or data frame.
- `sapply()`: to apply functions on a list/ vector/ data frame and returns an array or matrix object.
- `tapply()`: to compute functions to each factor variable in a vector.
- `lapply()`: to apply functions on list/ vector/ matrix objects and returns a list object.

Examples

1. Let consider two variables: salaries and jobs. Compute the mean for each level of jobs.

```
1 > salaries <- c(80000, 62000, 113000, 68000, 75000, 79000,
2   112000, 118000, 65000, 117000)
3 > jobs <- c('DS', 'DA', 'DE', 'DA', 'DS', 'DS', 'DE', 'DE',
4   'DA', 'DE') # factor variable
5 > print(tapply(salaries, jobs, mean))
      DA      DE      DS
65000 115000  78000
```

2. Create a new function add_one that, given an input, adds 1 unit. Then, define a vector my_vector and print a vector by applying add_one at each element.

```
1 > add_one <- function(x) x+1
2 > my_vector <- c(1, 2, 3)
3 > print(sapply(my_vector, add_one))
4 [1] 2 3 4
```